

Fig A1 - which model performs best? (matched at 8M forward passes)

single DeepChem scaffold hold-out (rarest scaffolds in test) · error bars = ± 1 sd over pretraining seeds for 6 arms, over head seeds only for: no_pretrain (frozen), unsup→sup (MLM→SFT) · classical anchors: 3 XGBoost seeds · no_pretrain = random-init encoder, frozen

