

Fig A1 (H1) — headline comparison at 8M forward passes (single pretraining seed)

no_pretrain = random encoder, frozen, head on eval · no_pretrain_end_to_end (E1) = random encoder finetuned end-to-end · sup_only = random init → SFT (no MLM)

