Fig A1.a - which model performs best? (matched at 8M forward passes)

single DeepChem scaffold hold-out (rarest scaffolds in test; 113-204 test molecules on ESOL/BACE/BBBP) · error bars = ±1 sd over pretraining seeds where replicates exist, else over head seeds (7 arms have pretraining-seed replicates) · classical anchors: 3 XGBoost seeds · no_pretrain (frozen) = random-init encoder, features frozen; no_pretrain (end-to-end) = same encoder fine-tuned on the task