


Fig A1.b - which model performs best? (matched at 8M forward passes)

pooled 5-fold scaffold cross-validation (README §8.1 protocol of record) · error bars = ± 1 sd across the 5 scaffold folds
 · classical anchors: 3 XGBoost seeds · no_pretrain (frozen) = random-init encoder, features frozen; no_pretrain (end-to-end) = same encoder fine-tuned on the task



Morgan+XGBoost (ECFP4) no_pretrain (end-to-end) sup_only: sparse_all unsup→sup: dense (MTR)
 Morgan+desc+XGBoost unsup_only (MLM) sup_only: dense+sparse unsup→sup: dense+sparse
 no_pretrain (frozen) sup_only: dense (MTR) unsup→sup: sparse_all random (chance)