

Fig A1.b - which model performs best? (matched at 8M forward passes)

pooled 5-fold scaffold cross-validation (README §8.1 protocol of record) · error bars = ± 1 sd across the 5 scaffold folds
 · classical anchors: 3 XGBoost seeds · no_pretrain (frozen) = random-init encoder, features frozen; no_pretrain (end-to-end) = same encoder fine-tuned on the task

