


Fig A2.b - scaling of the primary regimes in unique molecules seen

pooled 5-fold scaffold CV · error bars = ± 1 sd across the 5 folds (no head-seed, no pretraining-seed spread), so every interval means the same thing · the classical anchors and no_pretrain baselines are omitted for legibility — they are bars in A1.a/A1.b · x = unique molecules, so a ladder that runs VERTICAL is spending compute on repetition, not on new chemistry; the sparse SFT pool caps at ~0.52M and the unsupervised corpus at 12M



● unsup_only (MLM) ▲ sup_only: dense (MTR) ■ sup_only: dense+sparse
 ▲ unsup→sup: dense (MTR) ▼ sup_only: sparse_all corpus exhausted (12M molecules)