

Fig B1p1 - label-efficiency and mechanism: does the frozen probe fit, or generalize?

frozen probe re-fit at each label budget on the 8M encoders · both SFT arms are the SAME 'dense' (RDKit-MTR) recipe, so sup_only vs unsup→sup differs only in the MLM stage -- neither is an average over recipes · 3 subsample draws × 3 head seeds per point · HIV scored by ROC-AUC here so train and test share a metric

