


Fig B1p1 - label-efficiency and mechanism: does the frozen probe fit, or generalize?

frozen probe re-fit at each label budget on the 8M encoders · both SFT arms are the SAME 'dense' (RDKit-MTR) recipe,
so sup_only vs unsup→sup differs only in the MLM stage -- neither is an average over recipes · 3 subsample draws × 3
head seeds per point · HIV scored by ROC-AUC here so train and test share a metric



no_pretrain (frozen) sup_only: dense (MTR) unsup→sup: dense (MTR) no_pretrain (end-to-end) test train
unsup_only (MLM, no labels)