


Fig B2 - does content-free pretraining help just as much?

zero = no_pretrain (end-to-end), i.e. skipping pretraining and just fine-tuning a random encoder · pooled 5-fold scaffold CV. Each corrupted arm is matched to its real counterpart in objective, data volume, compute and schedule; only chemical content is destroyed.

lift over no_pretrain (end-to-end) (%)



■ unsup_only (MLM) - real ■ corrupted MLM (shuffled tokens)
■ sup_only: dense (MTR) - real ■ corrupted MTR (shuffled targets)