


Fig C1J1 - supervised pretraining data: which type helps, how much, and does it follow chemistry?



deduped wave: the SFT blacklist drops 34,301 eval molecules, so no arm trains on eval-test molecules. Arms are unsup→sup from one shared 2M-FP MLM base, all lifted against no_pretrain (end-to-end) → the random encoder FINE-TUNED on the task, not the frozen one, since clearing a frozen random encoder is close to automatic. SCHEME: pooled 5-fold scaffold CV. (c) omits dense (MTR) and dense+sparse: descriptor regression has no family molecule set to measure similarity against. Families are sampled, so similarities are lower bounds.