



Fig C1J1 - supervised pretraining data: which type helps, how much, and does it follow chemistry?



deduped wave: the SFT blacklist drops 34,301 eval molecules, so no arm trains on eval-test molecules. Arms are unsup→sup from one shared 2M-FP MLM base, each lifted against the random-init floor scored in its own wave.

(c) omits dense (MTR) and dense+sparse: descriptor regression has no family molecule set to measure similarity against. Families are sampled, so similarities are lower bounds.