



Which SFT pretrain helps? Which type helps, how much, and does it follow chemistry?


a (all arms: unsup→sup from one 2M-FP MLM base)



b SFT family → eval task



c H10 test: lift vs chemical similarity (n=24, Pearson $r=-0.14$, Spearman $\rho=-0.10$, $p=0.64$)



deduped wave: the SFT blocklist drops 34,301 eval molecules, so no arm trains on eval-test molecules. Arms are unsup→sup from one shared 2M-FP MLM base, each lifted against the random-init floor scored in its own wave.

SCHEME: single DeepChem scaffold hold-out (this wave has no 5-fold CV eval), so these lifts are NOT comparable with the CV numbers in A1.b / A2.

(c) omits dense (MTR) and dense+sparse: descriptor regression has no family molecule set to measure similarity against. Families are sampled, so similarities are lower bounds.